


Chapter 4

Greedy Algorithms




Slides by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.


4.5 Minimum Spanning Tree

Minimum Spanning Tree

Minimum spanning tree. Given a connected graph $G = (V, E)$ with real-valued edge weights c_e , an MST is a subset of the edges $T \subseteq E$ such that T is a spanning tree whose sum of edge weights is minimized.



$G = (V, E)$



$T, \sum_{e \in T} c_e = 50$

Cayley's Theorem. There are n^{n-2} spanning trees of K_n .

↑
can't solve by brute force

Applications

MST is fundamental problem with diverse applications.

- Network design.
 - telephone, electrical, hydraulic, TV cable, computer, road
- Approximation algorithms for NP-hard problems.
 - traveling salesperson problem, Steiner tree
- Indirect applications.
 - max bottleneck paths
 - LDPC codes for error correction
 - image registration with Renyi entropy
 - learning salient features for real-time face verification
 - reducing data storage in sequencing amino acids in a protein
 - model locality of particle interactions in turbulent fluid flows
 - autoconfig protocol for Ethernet bridging to avoid cycles in a network
- Cluster analysis.

Greedy Algorithms

Kruskal's algorithm. Start with $T = \emptyset$. Consider edges in ascending order of cost. Insert edge e in T unless doing so would create a cycle.


Reverse-Delete algorithm. Start with $T = E$. Consider edges in descending order of cost. Delete edge e from T unless doing so would disconnect T .

Prim's algorithm. Start with some root node s and greedily grow a tree T from s outward. At each step, add the cheapest edge e to T that has exactly one endpoint in T .

Remark. All three algorithms produce an MST.

Cycles and Cuts

Cut. A cut for a graph $G=(V,E)$ is a subset of nodes S .




Cut $S = \{4, 5, 8\}$
Crossing edges = 5-6, 5-7, 3-4, 3-5, 7-8

- An edge e **cross** a cut S if e has an endpoint in S and the other one in $V-S$.
- A cut S is crossed by X if some edge of X crosses S

Greedy Algorithms

Cut property. Suppose X is included in the set of edges of a MST for $G=(V,E)$. Pick a cut S that is not crossed by X . If e is the lightest edge that crosses S then $X \cup e$ is also part of a MST.




X = red edges
 e is in the MST

Greedy Algorithms

Cut property. Suppose X is included in the set of edges of a MST for $G=(V,E)$. Pick a cut S that is not crossed by X . If e is the lightest edge that crosses S then $X \cup e$ is also part of a MST.

Pf. (exchange argument)


- Let $T^*=(V,E^*)$ be a MST for G such that $X \subseteq E^*$
- Suppose e does not belong to T^* , and let's see what happens.
- Adding e to T^* creates a cycle C in $T^* \cup e$ that contains e .
- There exists another edge, say f , that is in C and cross S .
- f does not belong to X because X does not cross S
- $T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree.
- Since $c_e \leq c_f$, $\text{cost}(T') \leq \text{cost}(T^*)$.
- T' is also a MST and $X \cup e$ is part of it. ■



Prim's Algorithm

Prim's algorithm. [Jarník 1930, Dijkstra 1957, Prim 1959]

- Initialize S = any node.
- Apply cut property to S .
- Add min cost edge that crosses S to T , and add one new explored node u to S .



Bad Implementation: Prim's Algorithm

Implementation (Naïve)

- Maintain set of explored nodes S .
- Find the lightest edge that crosses S in $O(m)$ time
- Total complexity $O(m.n)$

Good Implementation: Prim's Algorithm

Implementation. Use a priority queue.


- Maintain set of explored nodes S .
- For each unexplored node v , maintain attachment cost $a[v]$ = cost of cheapest edge v to a node in S .
- $O(n^2)$ with an array; $O(m \log n)$ with a binary heap.

```
Prim(G, c) {  
    foreach (v ∈ V) a[v] ← ∞  
    Initialize an empty priority queue Q  
    foreach (v ∈ V) insert v onto Q  
    Initialize set of explored nodes S ← ∅  
  
    while (Q is not empty) {  
        u ← delete min element from Q  
        S ← S ∪ { u }  
        foreach (edge e = (u, v) incident to u)  
            if ((v ∉ S) and (ce < a[v]))  
                decrease priority a[v] to ce  
    }  
}
```


Kruskal's Algorithm

Kruskal's algorithm. [Kruskal, 1956]

- Consider edges in ascending order of weight.
- Case 1: If adding e to T creates a cycle, discard e
- Case 2: Otherwise, insert $e = (u, v)$ into T



Case 1




Case 2


Kruskal's Algorithm: Proof of correctness

Kruskal's algorithm. [Kruskal, 1956]

- Case 1: If adding e to T creates a cycle, discard e
 - Optimal solution does not have a cycle
- Case 2: Otherwise, insert $e = (u, v)$ into T
 - Pick the cut S as the nodes that are reachable from u in T



Case 1



Case 2

Kruskal's Algorithm: Bad Implementation

Kruskal's algorithm. [Kruskal, 1956]

- Sorting the edges $O(m \log m)$
- Testing the existence of a cycle while considering edge e : $O(n)$ via a DFS(BFS). Note that a tree has at most n edges.
- For all edges $O(m.n)$
- Total complexity $O(m \log m) + O(m n) = O(n.m)$

Implementation: Kruskal's Algorithm

Implementation. Use the **union-find** data structure.

- Build set T of edges in the MST.
- Maintain set for each connected component.
- $O(m \log n)$ for sorting and $O(m \underbrace{\alpha(m, n)}_{\text{essentially a constant}})$ for union-find.

\swarrow $m \leq n^2 \Rightarrow \log m$ is $O(\log n)$ \searrow essentially a constant


```
Kruskal(G, c) {  
    Sort edges weights so that  $c_1 \leq c_2 \leq \dots \leq c_m$ .  
     $T \leftarrow \phi$   
  
    foreach ( $u \in V$ ) make a set containing singleton  $u$   
  
    for  $i = 1$  to  $m$     are  $u$  and  $v$  in different connected components?  
         $(u, v) = e_i$   $\swarrow$   
        if ( $u$  and  $v$  are in different sets) {  
             $T \leftarrow T \cup \{e_i\}$   
            merge the sets containing  $u$  and  $v$   
        }  $\nwarrow$  merge two components  
    return  $T$   
}
```

Greedy Algorithms


Cycle property. Let C be any cycle in G , and let f be an edge belonging to C . If $\text{cost}(f) > \text{cost}(e)$ for every e in C , then f cannot belong to a MST.

Pf. (exchange argument)

- Suppose f belongs to a MST T^* , and let's see what happens.
- Deleting f from T^* creates two connected components.
- There is at least one edge in $C-f$, say e , that connects the two components. Thus,
- $T' = T^* \cup \{e\} - \{f\}$ is also a spanning tree.
- Since $c_e < c_f$, $\text{cost}(T') < \text{cost}(T^*)$.
- This is a contradiction. ■



4.7 Clustering



Outbreak of cholera deaths in London in 1850s.
Reference: Nina Mishra, HP Labs

Clustering

Clustering. Given a set U of n objects labeled p_1, \dots, p_n , classify into coherent groups.

↑
photos, documents, micro-organisms

Distance function. Numeric value specifying "closeness" of two objects.

↑
number of corresponding pixels whose intensities differ by some threshold

Fundamental problem. Divide into clusters so that points in different clusters are far apart.

- Routing in mobile ad hoc networks.
- Identify patterns in gene expression.
- Document categorization for web search.
- Similarity searching in medical image databases
- Skycat: cluster 10^9 sky objects into stars, quasars, galaxies.

Clustering of Maximum Spacing


k-clustering. Divide objects into k non-empty groups.

Distance function. Assume it satisfies several natural properties.


- $d(p_i, p_j) = 0$ iff $p_i = p_j$ (identity of indiscernibles)
- $d(p_i, p_j) \geq 0$ (nonnegativity)
- $d(p_i, p_j) = d(p_j, p_i)$ (symmetry)

Spacing. Min distance between any pair of points in different clusters.

Clustering of maximum spacing. Given an integer k , find a k -clustering of maximum spacing.



Clustering of Maximum Spacing



$K=3$

Spacing = 2

Greedy Clustering Algorithm

Single-link k-clustering algorithm.

- Form a graph on the vertex set U , corresponding to n clusters.
- Find the closest pair of objects such that each object is in a different cluster, and add an edge between them.
- Repeat $n-k$ times until there are exactly k clusters.

Key observation. This procedure is precisely Kruskal's algorithm (except we stop when there are k connected components).

Remark. Equivalent to finding an MST and deleting the $k-1$ most expensive edges.

Greedy Clustering Algorithm: Analysis

Lemma. If $C = \{C_1, \dots, C_k\}$ and $C' = \{C'_1, \dots, C'_k\}$ are two different clustering then there exists a pair of points p_i, p_j such that p_i and p_j are in the same cluster in C and in different clusters in C'

Pf. Homework


.

Greedy Clustering Algorithm: Analysis

Theorem. Let C^* denote the clustering C^*_1, \dots, C^*_k formed by the single k -link algorithm. C^* is a k -clustering of max spacing.

Pf. Let C denote some other clustering C_1, \dots, C_k .

- The spacing of C^* is the length d^* of lightest edge that connects two cluster of C^* .
- Let p_i, p_j be in the same cluster in C^* , say C^*_r , but different clusters in C , say C_s and C_t .
- Some edge (p, q) on p_i - p_j path in C^*_r connects two different clusters in C .
- All edges on p_i - p_j path have length $\leq d^*$ since Kruskal chose them.
- Spacing of C is $\leq d^*$ since p and q are in different clusters. ■



Extra Slides

MST Algorithms: Theory

Deterministic comparison based algorithms.

- $O(m \log n)$ [Jarník, Prim, Dijkstra, Kruskal, Boruvka]
- $O(m \log \log n)$. [Cheriton-Tarjan 1976, Yao 1975]
- $O(m \beta(m, n))$. [Fredman-Tarjan 1987]
- $O(m \log \beta(m, n))$. [Gabow-Galil-Spencer-Tarjan 1986]
- $O(m \alpha(m, n))$. [Chazelle 2000]

Holy grail. $O(m)$.

Notable.

- $O(m)$ randomized. [Karger-Klein-Tarjan 1995]
- $O(m)$ verification. [Dixon-Rauch-Tarjan 1992]


Euclidean.

- 2-d: $O(n \log n)$. compute MST of edges in Delaunay
- k-d: $O(k n^2)$. dense Prim

Dendrogram


Dendrogram. Scientific visualization of hypothetical sequence of evolutionary events.

- Leaves = genes.
- Internal nodes = hypothetical ancestors.



Dendrogram of Cancers in Human

Tumors in similar tissues cluster together.



Reference: Botstein & Brown group


■ gene expressed
■ gene not expressed

Union-Find on disjoint sets

Motivation

- Perform repeated union and find operations on disjoint data sets.
- Examples:
 - Kruskal's MST algorithm
 - Connected Components
- Goal: define an ADT that supports Union-Find queries on disjoint data sets efficiently.

Example: connected components



- *Initial set* $S = \{\{a\}, \{b\}, \{c\}, \{d\}, \{e\}, \{f\}, \{g\}, \{h\}, \{i\}, \{j\}\}$
- (b,d) $S = \{\{a\}, \{b,d\}, \{c\}, \{e\}, \{f\}, \{g\}, \{h\}, \{i\}, \{j\}\}$
- (e,g) $S = \{\{a\}, \{b,d\}, \{c\}, \{e,g\}, \{f\}, \{h\}, \{i\}, \{j\}\}$
- (a,c) $S = \{\{a,c\}, \{b,d\}, \{e,g\}, \{f\}, \{h\}, \{i\}, \{j\}\}$
- (h,i) $S = \{\{a,c\}, \{b,d\}, \{e,g\}, \{f\}, \{h,i\}, \{j\}\}$
- (a,b) $S = \{\{a,c,b,d\}, \{e,g\}, \{f\}, \{h,i\}, \{j\}\}$
- (e,f) $S = \{\{a,c,b,d\}, \{e,f,g\}, \{h,i\}, \{j\}\}$
- (b,c) $S = \{\{a,c,b,d\}, \{e,f,g\}, \{h,i\}, \{j\}\}$

Union-Find Abstract Data Type


- Let $S = \{S_1, S_2, \dots, S_k\}$ be a dynamic collection of disjoint sets.
- Each set S_i is identified by a representative member.
- Operations:
 - Make-Set(x): create a new set S_x , whose only member is x (assuming x is not already in one of the sets).
 - Union(x, y): replace two disjoint sets S_x and S_y represented by x and y by their union.
 - Find-Set(x): find and return the representative of the set S_x that contains x .

Disjoint sets: linked list representation


- Each set is a linked list, and the representative is the head of the list. Elements point to the successor and to the head of the list.
- Make-Set: create a new list: $O(1)$.
- Find-Set: search for an element down the list: $O(1)$.
- Union: link the tail of L_1 to the head of L_2 , and make each element of L_2 point to the head of L_1 : $O(|L_2|)$.
- A sequence of n Make-Set operations $+(n-1)$ Union operations may take $n + \sum i = O(n^2)$ operations.

Example: linked list representation

$S_1 = \{c, h, e, b\}$




$S_2 = \{f, g, d\}$



Example: linked list representation

$$S_1 \cup S_2 = \{c, h, e, b\} \cup \{f, g, d\}$$



Weighted-union heuristic

- When doing the union of two lists, append the shorter one to the longer one.
- A single operation will still take $O(n)$ time.
 - Two lists of $n/2$ elements

Weighted-union heuristic



- However, for a sequence of $m > n$ Make-Set, Union, and Find-Set operations, of which n are Make-Set, the total time is $O(m + n \lg n)$ instead of $O(mn)$!
- Proof outline:
 - The total of operations equals the number of times an the objects have its pointers updates
 - Whenever an object x has its pointer updated, the size of the list where x lies doubles. Thus, an object x can only have its pointer updated $\log n$ times

Disjoint sets: tree representation

- Each set is a tree, and the representative is the root.
- Each element points to its parent in the tree. The root points to itself.
- Make-Set: takes $O(1)$.
- Find-Set: takes $O(h)$ where h is the height of the tree.
- Union is performed by finding the two roots, and choosing one of the roots, to point to the other. This takes $O(h)$.
- The complexity therefore depends on how the trees are maintained!

Example

$S = \{S_1, S_2\}$, $S_1 = \{c, h, e, b\}$, $S_2 = \{f, g, d\}$



Union by rank

- We want to make the trees as shallow as possible → trees must be balanced.
- When taking the union of two trees, make the root of the shorter tree become the child of the root of the longer tree.
- Keep track of height of each sub-tree: → keep the *rank* of each node.
- Every time Union is performed, update the rank of the root.

Complexity of Find-Set (1)

- Claim: A tree of height h has at least 2^h nodes
- Proof: Assume that the property holds before the k -th operation and prove that it is true after the k -th operation.

Complexity of Find-Set (2)

- Case 1) **The tree height does not grow.** It follows that one tree was shorter than the other, in which case it is clearly true, because h didn't grow and the number of nodes did.
- Case 2) **The height does grow to $h+1$.** It follows that both tree heights were the same. By hypothesis, each sub-tree has at least 2^h nodes, so the new tree has at least $2 \cdot 2^h = 2^{h+1}$ nodes. Thus, the height of the tree grew by 1 to $h+1$, which proves the induction step.

Tree representation


- For a sequence of $m > n$ Make-Set, Union, and Find-Set operations, of which n are Make-Set, the total time is $O(m \lg n)$.

Path compression


- Speed up Union-Find operations by shortening the sub-tree paths to the root.
- During a Find-Set operation, make each node on the find path point directly to the root.
- worst-case time complexity is $O(m \alpha(n))$ where $\alpha(n)$ is the *inverse Ackerman function*.
- The inverse Ackerman function grows so slowly that for all practical purposes $\alpha(n) \leq 4$ for very large n .

Example: path compression

Before



After



Summary


- Union-Find has many applications.
- For a sequence of $m > n$ Make-Set, Union, and Find-Set operations, of which n are Make-Set:
 - List implementation: $O(m + n \lg n)$ with weighted union heuristic.
 - Tree implementation: union by rank + path compression yields $O(m \alpha(n))$ complexity.

Vertex Cover

VERTEX COVER: Given a graph $G = (V, E)$ and an integer k , is there a subset of vertices $S \subseteq V$ such that $|S| \leq k$, and for each edge, at least one of its endpoints is in S ?

Ex. Is there a vertex cover of size ≤ 4 ? Yes.

Ex. Is there a vertex cover of size ≤ 3 ? No.



● vertex cover

Vertex Cover

- . Consider the Greedy Algorithm that always selects the node with minimum degree?
- . Does it always produce the optimal solution? Why?
- . How can we implement it efficiently?